

RoundFi — Overview

Turning the savings circles people already trust into portable, on-chain credit

RoundFi

2026-06-23

CONTENTS

1. The thesis, in one sentence
2. The problem: most of the world is credit-invisible
3. The solution: three layers
4. The feedback loop
5. It's live: proof, not a deck
6. What we are honest about

Derived document. This is a derivation of [docs/spec/MASTER-SPEC.md](#) (§1 One sentence, §2 Problem, §3 Solution, §11 Deployed state). The Master Spec is the single source of truth; if a number here disagrees with it, the Master Spec wins. Every claim of deployment is pinned to the validated devnet run of 2026-06-12 ([docs/operations/v52-devnet-run-book.md](#)).

1. The thesis, in one sentence

RoundFi is infrastructure that turns verifiable financial behavior into portable, on-chain reputation.

Underneath the product is a **rotating savings-and-credit association** — the same instrument known as a *consórcio* or *junta* in Brazil, a *tanda* in Mexico, a *susu* in West Africa, a *chama* in Kenya — implemented on Solana. But the ROSCA is not the product. It is the **data engine**. Every cycle of a pool produces a stream of verifiable, financially-consequential events: a member paid on time, paid late, completed the round, or defaulted. The **reputation** built from that stream is the product — a portable credit signal a wallet carries from pool to pool and, eventually, from protocol to protocol.

The distinction matters because it inverts the usual pitch. RoundFi is **not** "a savings app with a score bolted on." The savings mechanism exists to **generate behavior worth scoring**; the score exists to make future credit cheaper for people who have proven they pay. That is the entire business.

2. The problem: most of the world is credit-invisible

Thin-file and no-file individuals — the majority of adults across Brazil, Latin America, Africa, and Southeast Asia — have **no portable record** of whether they keep financial commitments. The crucial point, and the one most lending products get wrong, is this: **they are not high-risk; they are unmeasured**. It is the *absence of a track record*, not the presence of bad behavior, that excludes them from affordable credit.

These same populations are not financially passive. They already run informal credit instruments at scale — ROSCAs, *consórcios*, *tandas*, *susus*, *chamas* — and those instruments generate **exactly the behavior a lender would want to see**: recurring contributions, social enforcement of repayment, and completion of multi-month obligations.

The problem is that the behavior **evaporates**. It lives in a WhatsApp group or a paper notebook. It does not survive the group dissolving. And it cannot be presented to anyone outside the group — not to a bank, not to a lender, not to a future pool of strangers. A person can complete a dozen honest savings circles and arrive at a credit application with **nothing to show for it**.

The gap RoundFi closes is precise: take an instrument people **already trust and use**, run it on rails that make every event **verifiable and durable**, and emit a **portable reputation** from it.

3. The solution: three layers

RoundFi is built as three layers, each feeding the next.

3.1 Layer one — the pool (a ROSCA on-chain)

A fixed set of members each contribute a recurring installment. Each cycle, one member receives the pooled **carta** (the credit draw). The pool is not a trust exercise: it is protected by a layered defense the spec calls the **Triple Shield** — a member **stake** locked at join, an **escrow** slice taken from each contribution, and a **solidarity reserve** ("Cofre Solidário") that covers the first tranche of any default. Every contribution, every payout, and every default is a Solana transaction.

Crucially, the pool is **self-funding**. A viability guard checks that the float retained from contributions is always enough to pay the drawn member every cycle, with no external subsidy — RoundFi refuses to create a pool whose math does not close. (The mechanics live in [02-technical-whitepaper](#) and the solvency math in [05-stress-lab-economic-model](#).)

3.2 Layer two — the attestation layer

Each consequential event emits a signed, **immutable on-chain attestation**: a behavioral record that says *paid*, *late*, *pool-completed*, *defaulted*, or *payout-claimed*. Attestations are keyed by **subject = wallet**. This is the verifiable substrate — the part that does **not** evaporate. Where a notebook forgets and a WhatsApp group dissolves, an attestation is a permanent, auditable fact.

3.3 Layer three — the reputation engine

An off-chain scorer and an on-chain level ladder consume the attestation stream and produce three things: a **score**, a **tier** (L1–L4), and four **behavioral metrics** (Reliability, Punctuality, Commitment, Recovery). The tier feeds back into the protocol — a higher tier means a **lower required stake** — and the reputation is exposed to external consumers (a lender, a fintech) via an API. The full engine is documented in [04-behavioral-reputation-score](#).

3.4 The design principle that holds it together

One separation runs through the entire protocol: **receiving capital is not merit; keeping your obligations after receiving it is**. This is enforced on-chain, not by convention. *Claiming* your carta is deliberately **score-neutral** — the credit for completing a pool only lands once a member has paid through to the **end**, after they already received their draw. You cannot take the money and farm a reputation; the reputation is earned by paying it *back*.

4. The feedback loop

The loop is the whole point, and it is short enough to state in one line:

Good behavior in the pool lowers your cost of capital in the next pool — and the resulting reputation is yours to carry.

A member who climbs the ladder from L1 to L4 (Elite) goes from locking **50%** of the credit as stake down to just **3%** — roughly a 94% reduction in the capital they must put up. That discount is a direct, legible reward for a proven track record, and the protocol's risk does **not** rise as the stake falls, because the upper tiers are gated on a long history of completed pools that is itself the risk signal.

The endgame extends past RoundFi's own pools. A wallet's portable output —

`{ tier, reliability, punctuality }` — is exactly what an external lender (Kamino, MarginFi, a fintech) can consume to **price credit for a borrower who would otherwise be invisible**. The savings circle becomes the on-ramp to the formal credit system.

5. It's live: proof, not a deck

This is the part a deck cannot fake. RoundFi is **not** a roadmap of intentions — the protocol is deployed and was exercised **end-to-end on devnet with real USDC** on 2026-06-12. **Seven of the eight** protocol capability areas were run through to completion; the eighth (default settlement) is armed and time-gated, waiting only on its grace window to elapse.

#	Capability	Status
1	Full pool lifecycle + Pass-3 reputation scoring	
2	<code>close_pool</code> (terminal Closed state)	
3	Yield Cascade (init → deposit → harvest)	
4	Escape Valve — direct (list → buy)	
5	Pause circuit-breaker (pause → gate → unpause)	
6	Escrow vesting (on-time stake release)	
7	Escape Valve — commit-reveal (anti-MEV)	
8	Rent-reclaim ceremony (full teardown)	
9	Default settlement	armed; grace-gated

The rent-reclaim ceremony (#8) is the **true** end of a pool's life — not just flipping status to `Closed`, but reclaiming every lamport of rent back to members and the authority. On the validated run, the authority's SOL balance went **up** (+0.0108 net of fees): the protocol cleans up after itself.

The reputation separation that is RoundFi's thesis was verified on-chain in the same run (pool `Ga2RwgSk...`): the *payment*, *payout-claimed*, and *pool-complete* events were emitted **distinctly**, with the payout-claimed event correctly landing in the **neutral** bucket — receiving capital scored nothing, completing the pool scored the reward.

The deployed programs (devnet, shared across clusters):

Program	ID
<code>roundfi-core</code>	<code>8LVrgxKwKwqjcdq7rUUwWY2zPNk8anpo2JsaR9jTQQjw</code>
<code>roundfi-reputation</code>	<code>Hpo174C6JTCfiZ6r8VYVQdKxo3LBHaJmMbkgREkxe9R2</code>
<code>roundfi-yield-mock</code>	<code>GPTMPGxexhwkhXNovnfrcSsmoWPUhedvKAQfTV2Ef5AQ</code>
<code>roundfi-yield-kamino</code>	<code>74izMa4WzLuHvtzDLdNzcygKe5fYwtD95EiWMuzhFdb</code>

Transaction signatures and the reproducible runbook live in `docs/operations/v52-devnet-runbook.md` — every claim above can be re-verified, not taken on trust.

6. What we are honest about

Credibility is also about what is *not* yet done. Two of the four behavioral metrics — **Commitment** and **Recovery** — currently ship as `null` in the score API. They are deliberately **not** fabricated: they will be calibrated from real default-and-recovery sequences produced by the canary before they are published. The on-chain Elite gate is **v1-provisional** (score + completed pools); the sharper metric-based Elite criteria live off-chain until the weights are proven. And the protocol's hardest gates to mainnet are **non-code**: regulatory review (LGPD, Bacen) handled by counsel, an external audit, and a multi-sig on the treasury.

We would rather ship a `null` we can defend than a number we cannot. That posture — **legible, auditable, re-verifiable** — is the same property that makes the reputation worth something to a lender in the first place.

Cross-references: protocol mechanics → [02-technical-whitepaper](#) ; program topology → [03-architecture-spec](#) ; the reputation engine in full → [04-behavioral-reputation-score](#) ; the devnet evidence → [11-devnet-status-and-proof](#) . The complete source of truth is MASTER-SPEC §1–3 and §11.